

Anatomy of an evaluator

How Arize runs that prompt against your traces



Data fields

WHAT IT IS

The prompt names the span attributes to read as placeholders, which the evaluator fills with real trace data at run time.

WHY IT MATTERS

Binds the judge to the actual input, output, and context of each span rather than letting it guess.



Eval model

WHAT IT IS

The model that actually runs the judgement on each span or trace, scoring it against your prompt.

WHY IT MATTERS

The flexibility lives here: swap in a cheaper or stronger judge without rewriting a single line of the prompt.



Output

WHAT IT IS

Each run returns a label, an optional score, and an explanation, landing on the span as structured eval attributes.

WHY IT MATTERS

Constrained, machine-parseable results are what let you aggregate, filter, alert, and track quality over time.