

Arrange, act, eval, assert

The unit-test loop you already know, with one extra step for fuzzy AI output



Arrange

THE STEP


Set up your inputs — the rows of your golden dataset, each a trusted input paired with its expected output. Every row becomes one parameterised case.



Act

THE STEP


Run the thing under test. Your AI app processes each row and produces an output — once per row, exactly like the body of a parameterised test.



Eval

THE STEP

The new step. Fuzzy output isn't assertable, so grade it first with a code check or a model-as-judge — turning a squishy sentence into a score, a label, a pass or fail.



Assert

THE STEP

Check the result — but gate on an aggregate, not every row. A passing percentage, not a perfect board. Drop below your threshold and CI blocks the PR.